



Threads



Threads concepts


Thread

- An independent and schedulable execution unit.
- A process can be divided into **two or more running threads**.
- A single thread of control(= a UNIX process)
 - Each process is doing **only one thing at a time**.
- Multiple threads of control in a process.
 - The process can do **more than one thing at a time**.
- Multithreading is possible on even uni-processor
 - by **time-division multiplexing**.

Threads concepts

🖥️ A typical example

- Apache web server



Program code & data is identical.
But, PC(Program Counter) and
stack are different.

Threads concepts

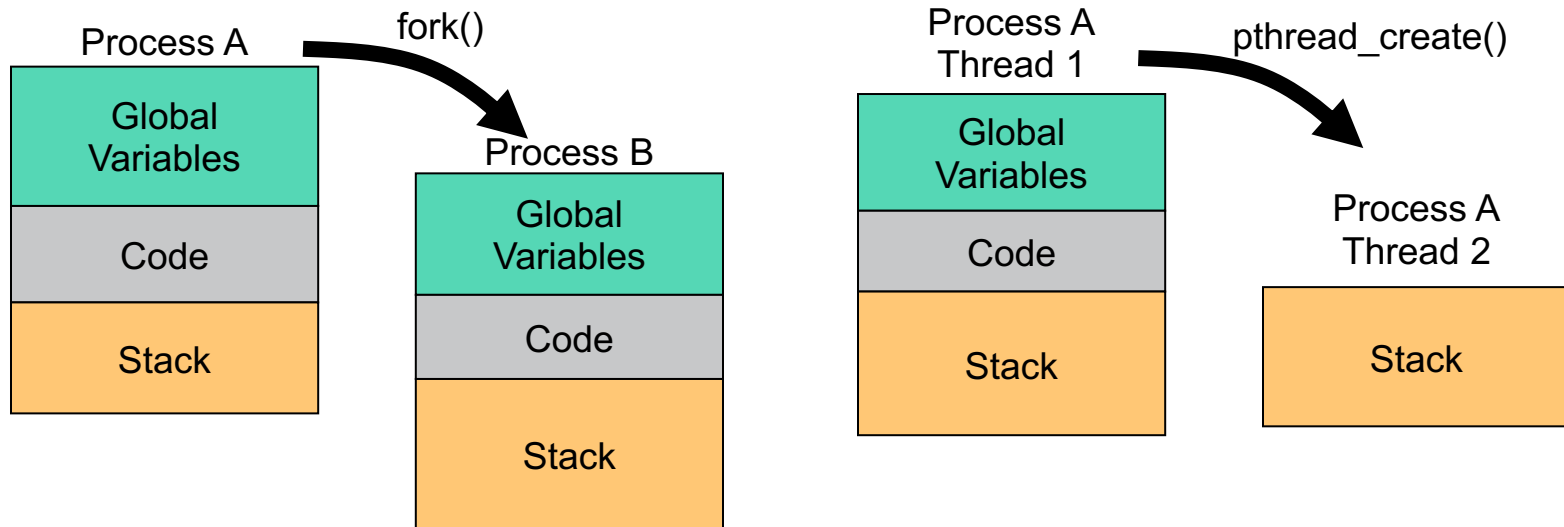
Advantages of thread

- Easy to share information.
 - the memory address space and file descriptors.
- Throughput can be improved.
 - The processing of independent tasks can be interleaved.
- More interactive.
 - The separated threads can deal with user input/output.

Threads concepts

Advantages of thread(cont.)

- The cost for creating a new process is low.



Threads concepts

A thread-specific information

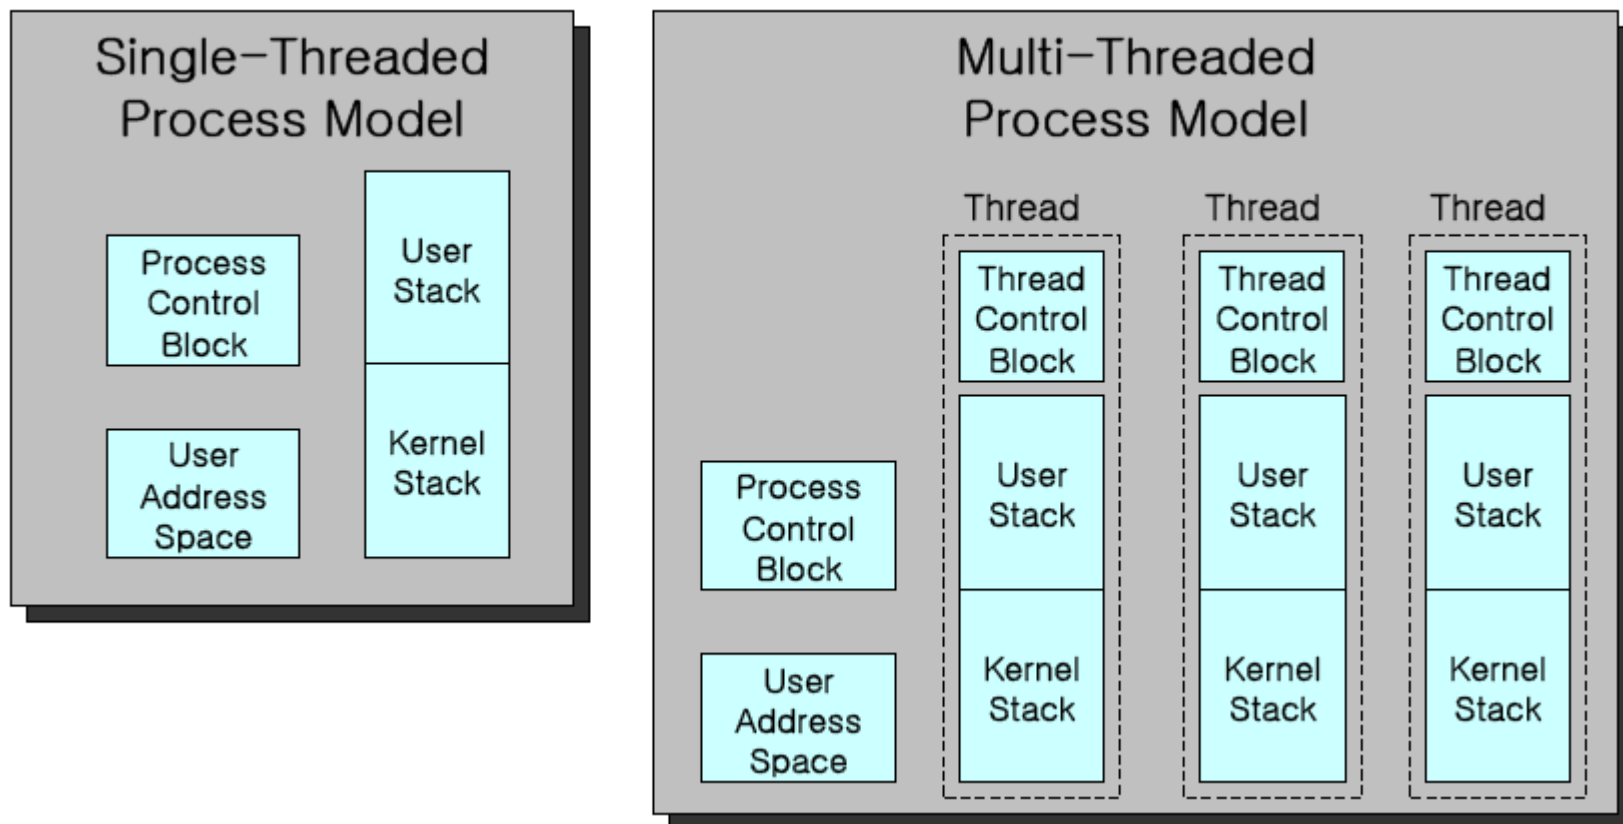
- Thread ID
- Register values
- Stack
- Scheduling priority
- A signal mask

Sharable information among threads in a process

- Text section
- Global data
- Heap
- File descriptor

Threads concepts

❏ Process vs. thread



Posix thread

What is pthread?

- IEEE POSIX 1003.1c standards

Pthread naming convention

- pthread_

Compiling pthread program

- `$ gcc -pthread xxx.c`

Thread identification

Thread ID

- Identifier of thread (similar to process ID.)
- A thread ID is represented by `pthread_t` data type.
 - Unsigned long integer in Linux.
 - A pointer to the `pthread` structure in FreeBSD.

Thread identification

```
#include <pthread.h>
```

```
pthread_t pthread_self(void);
```

Returns: the thread ID of the calling thread

 Obtain its own thread ID.

```
#include <pthread.h>
```

```
int pthread_equal(pthread_t tid1, pthread_t tid2);
```

Returns: nonzero if equal, 0 otherwise

 Compare two thread IDs

Thread creation

```
#include <pthread.h>
```

```
int pthread_create(pthread_t *tidp,  
                  const pthread_attr_t *attr,  
                  void *(*start_rtn)(void), void *arg);
```

Returns: 0 if OK, error number on failure

Create a new thread.

- *tidp* is the ID of the newly created thread.
- Execute *start_rtn* with *arg* as its argument.
- *attr* is set to NULL for the default attributes.

Thread creation

Example

```
#include "apue.h"
#include <pthread.h>

pthread_t ntid;

void printids(const char *s)
{
    pid_t    pid;
    pthread_t tid;

    pid = getpid();
    tid = pthread_self();
    printf("%s pid %u tid %u (0x%x)\n", s, (unsigned int)pid,
        (unsigned int)tid, (unsigned int)tid);
}
```

What if to use *ntid* instead of *tid*?

Thread creation

Example(cont.)

```
void *thr_fn(void *arg)
{
    printids("new thread: ");
    return((void *)0);
}

int main(void)
{
    int    err;

    err = pthread_create(&ntid, NULL, thr_fn, NULL);
    if (err != 0)
        err_quit("can't create thread: %s\n", strerror(err));
    printids("main thread:");
    sleep(1); // New thread can run before the old thread terminates.
    exit(0);
}
```

Thread creation

Execution

In Solaris

\$./a.out

When a thread is created, there is no guarantee which runs first.

main thread: pid 7225 tid 1 (0x1)

new thread: pid 7225 tid 4 (0x4)

\$

In FreeBSD

\$./a.out

FreeBSD uses a pointer to the thread data structure for its thread ID.

main thread: pid 14954 tid 134529024 (0x804c000)

new thread: pid 14954 tid 134530048 (0x804c400)

\$

In Linux

\$./a.out

main thread: pid 10043 tid 3480123136 (0xcf6e7700)

new thread: pid 10043 tid 3471726336 (0xceee5700)

\$

Thread termination

❏ If any thread within a process call `exit()`?

- The entire process terminates.


❏ A single thread can exit without terminating the entire process.

- The thread **can simply return** from the start routine.
- The thread can be **canceled by another thread** in the same process.
- The thread can call **`pthread_exit()`**.

Thread termination

```
#include <pthread.h>
```

```
void pthread_exit(void *rval_ptr);
```

 Terminates a calling thread.


- *rval_ptr* is available to other threads in the process calling the `pthread_join()`.

Thread termination

```
#include <pthread.h>
```

```
int pthread_join(pthread_t thread, void **rval_ptr);
```

Returns: 0 if OK, error number on failure

 Suspends execution of the calling thread until the target thread terminates.

- It is similar to **wait()**.
- *rval_ptr* argument
 - If the thread returned from start routine, it contains the return code.
 - If the thread was canceled, it is set to `PTHREAD_CANCELED`.
 - If we're not interested in a return value, it is set to `NULL`.

Thread termination

Example

```
#include "apue.h"
#include <pthread.h>

void *thr_fn1(void *arg)
{
    printf("thread 1 returning\n");
    return((void *)1);
}

void *thr_fn2(void *arg)
{
    printf("thread 2 exiting\n");
    pthread_exit((void *)2);
}
```

Thread termination

Example(cont.)

```
int main(void)
{
    int      err;
    pthread_t tid1, tid2;
    void      *tret;

    err = pthread_create(&tid1, NULL, thr_fn1, NULL);
    if (err != 0)
        err_quit("can't create thread 1: %s\n", strerror(err));

    err = pthread_create(&tid2, NULL, thr_fn2, NULL);
    if (err != 0)
        err_quit("can't create thread 2: %s\n", strerror(err));
```

Thread termination

Example(cont.)

```
err = pthread_join(tid1, &tret);
if (err != 0)
    err_quit("can't join with thread 1: %s\n", strerror(err));
printf("thread 1 exit code %d\n", (int)tret);

err = pthread_join(tid2, &tret);
if (err != 0)
    err_quit("can't join with thread 2: %s\n", strerror(err));
printf("thread 2 exit code %d\n", (int)tret);

exit(0);
}
```

Thread termination

execution

```
$ ./a.out  
thread 1 returning  
thread 2 exiting  
thread 1 exit code 1  
thread 2 exit code 2  
$
```

Thread termination

```
#include <pthread.h>
```

```
int pthread_cancel(pthread_t tid);
```

Returns: 0 if OK, error number on failure


 Cancel another thread in the same process.

- Cause the thread with *tid* to behave as if it had called `pthread_exit()`.
- It doesn't wait for the thread to terminate; it merely makes the request.

Thread termination

```
#include <pthread.h>
```

```
void pthread_cleanup_push(void (*rtn)(void *), void *arg);  
void pthread_cleanup_pop(int execute);
```

-  A thread can arrange for functions to be called when it exits. (like `atexit()`)
- The functions are called as thread cleanup handlers.
 - More than one cleanup handler can be established.
 - The handlers are recorded in a stack, so they are executed in the reverse order of the registrations.

Thread termination


Pthread_cleanup_push

- Pushes routine onto the top of the stack of cleanup handlers.
- *rtn*: cleanup handler function
- *arg*: a single argument

Pthread_cleanup_pop

- Pops the top cleanup handler from the current thread's cleanup handler stack.
- If *execute* is 0, the cleanup handler is not called; it just removes the cleanup handler on top of stack.

Thread termination

-  When the thread performs one of the followings, cleanup handlers are executed.
- Makes a call to `pthread_exit()`.
 - Responds to a cancellation request.
 - Invoke `pthread_cleanup_pop()` with a nonzero *execute* argument.

Thread termination

Example

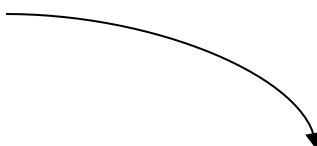
```
#include "apue.h"
#include <pthread.h>

void
cleanup(void *arg)
{
    printf("cleanup: %s\n", (char *)arg);
}
```

Thread termination

Example(cont.)

```
void *  
thr_fn1(void *arg)  
{  
    printf("thread 1 start\n");  
    pthread_cleanup_push(cleanup, "thread 1 first handler");  
    pthread_cleanup_push(cleanup, "thread 1 second handler");  
    printf("thread 1 push complete\n");  
  
    if (arg)  
        return((void *)1);  
  
    pthread_cleanup_pop(0);  
    pthread_cleanup_pop(0);  
    return((void *)1);  
}
```



We should match calls to `pthread_cleanup_pop` with the calls to `pthread_cleanup_push`; Otherwise, the program might not compile.

Thread termination

Example(cont.)

```
void *  
thr_fn2(void *arg)  
{  
    printf("thread 2 start\n");  
    pthread_cleanup_push(cleanup, "thread 2 first handler");  
    pthread_cleanup_push(cleanup, "thread 2 second handler");  
    printf("thread 2 push complete\n");  
  
    if (arg)  
        pthread_exit((void *)2);  
  
    pthread_cleanup_pop(0);  
    pthread_cleanup_pop(0);  
    pthread_exit((void *)2);  
}
```

Thread termination

Example(cont.)

```
int
main(void)
{
    int      err;
    pthread_t tid1, tid2;
    void      *tret;

    err = pthread_create(&tid1, NULL, thr_fn1, (void *)1);
    if (err != 0)
        err_quit("can't create thread 1: %s\n", strerror(err));

    err = pthread_create(&tid2, NULL, thr_fn2, (void *)1);
    if (err != 0)
        err_quit("can't create thread 2: %s\n", strerror(err));
```

Thread termination

Example(cont.)

```
err = pthread_join(tid1, &tret);
if (err != 0)
    err_quit("can't join with thread 1: %s\n", strerror(err));
printf("thread 1 exit code %d\n", (int)tret);

err = pthread_join(tid2, &tret);
if (err != 0)
    err_quit("can't join with thread 2: %s\n", strerror(err));
printf("thread 2 exit code %d\n", (int)tret);

exit(0);
}
```

Thread termination

execution

```
$ ./a.out  
thread 1 start  
thread 1 push complete  
thread 2 start  
thread 2 push complete  
cleanup: thread 2 second handler  
cleanup: thread 2 first handler  
thread 1 exit code 1  
thread 2 exit code 2  
$
```

Thread termination

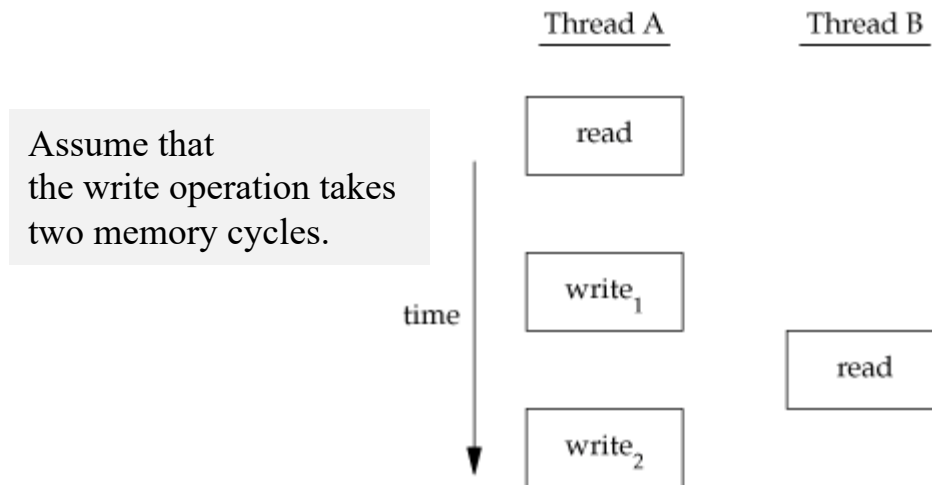
Comparison of process and thread primitives

Process primitive	Thread primitive	Description
fork	pthread_create	Create a new flow of control
exit	pthread_exit	Exit from an existing flow of control
waitpid	pthread_join	Get exit status from flow of control
atexit	pthread_cleanup_push	Register function to be called at exit from flow of control
getpid	pthread_self	Get ID for flow of control
abort	pthread_cancel	Request abnormal termination of flow of control

Thread synchronization

❏ Why thread synchronization?

- When one thread modify a variable that other threads read or modify, inconsistency problem exists.
- When multiple threads share the same memory, each thread must see a consistent view of its data.



Interleaved memory cycles with two threads.

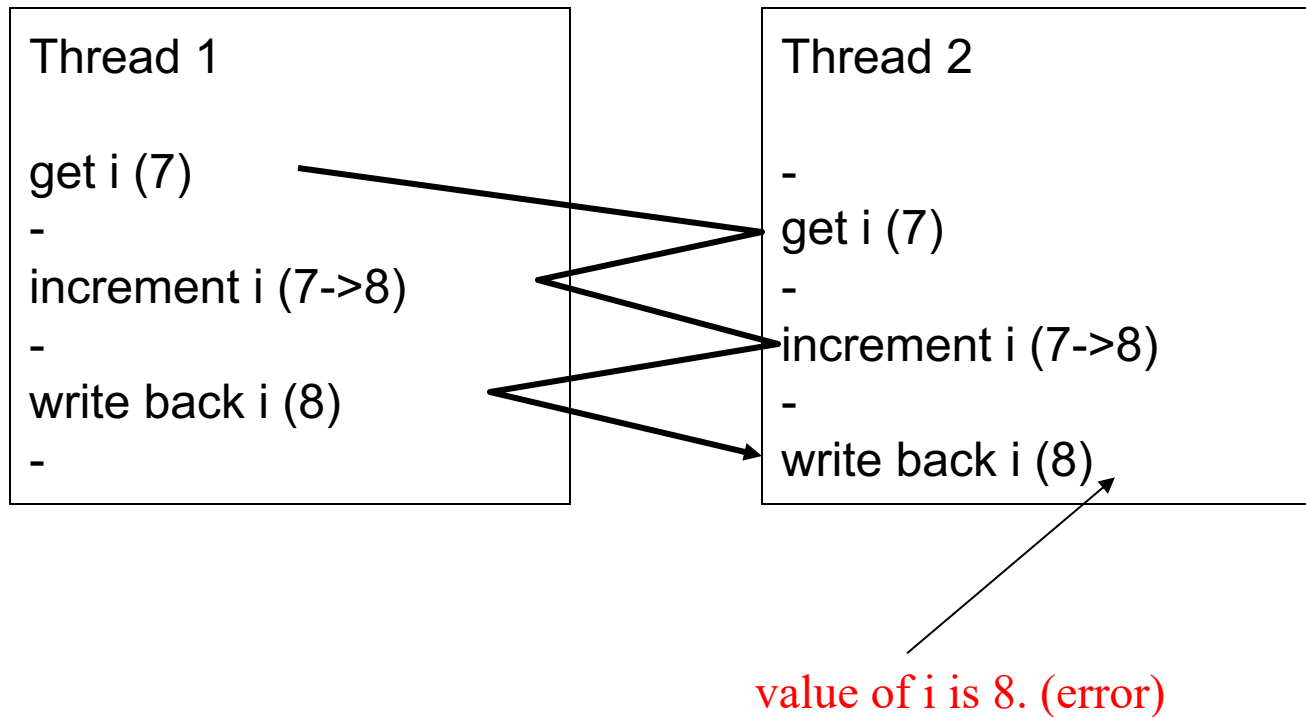
Thread synchronization

 Even a simple increment operation is broken down into three steps.

- Read the memory location into a register.
- Increment the value in the register.
- Write the new value back to the memory location.

Thread synchronization

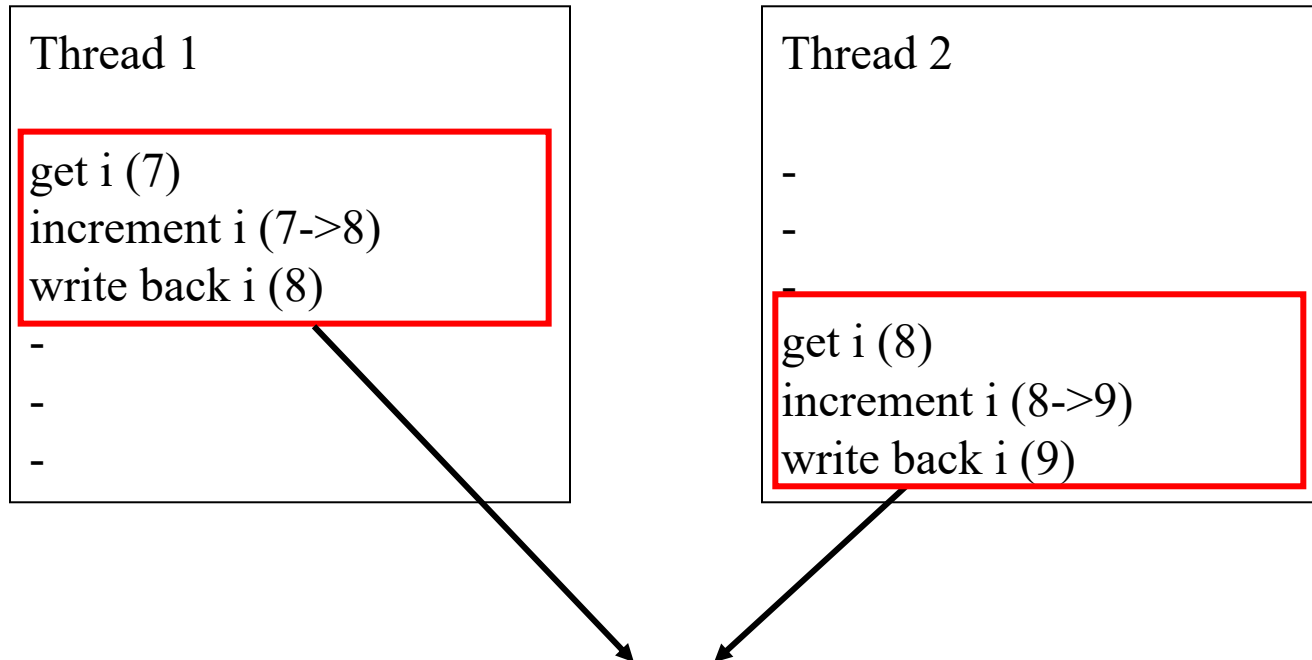
❏ If two threads try to increment the same variable at almost the same time?



Thread synchronization

❏ Synchronization is required.

- mutex, condition variable, ...



Mutexes

Mutex (mutual exclusion)

- A lock that we set before accessing a shared resource and release when we're done.
- While it is set, any other thread that tries to set it will block until it is released.
- When the mutex is released, all threads blocked on the lock will be unblocked.
- Finally, one of threads can set the lock, and the others will block again.

Mutexes

```
#include <pthread.h>
```

```
int pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *attr);  
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

Both return: 0 if OK, error number on failure

Pthread_mutex_init

- Mutex variable is represented by *pthread_mutex_t* data type.
- Before using a mutex, we must first initialize it.
- To initialize a mutex with the default attributes, *attr* is set to NULL.

Pthread_mutex_destroy

- Destroy the mutex.

Mutexes

```
#include <pthread.h>
```

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

All return: 0 if OK, error number on failure

Pthread_mutex_lock

- Lock a mutex.
- If the mutex is already locked, the calling thread will block until the mutex is unlocked.

Pthread_mutex_trylock

- Nonblocking version of pthread_mutex_lock().

Pthread_mutex_unlock

- Unlock a mutex.

Mutexes

Example

```
#include <stdlib.h>
#include <pthread.h>

struct foo {
    int      f_count;
    pthread_mutex_t f_lock;
    /* ... more stuff here ... */
};
```


Mutexes

Example(cont.)

```
struct foo *foo_alloc(void)                /* allocate the object */
{
    struct foo *fp;

    if ((fp = malloc(sizeof(struct foo))) != NULL) {
        fp->f_count = 1;
        if (pthread_mutex_init(&fp->f_lock, NULL) != 0) {
            free(fp);
            return(NULL);
        }
        /* ... continue initialization ... */
    }
    return(fp);
}
```

Mutexes

Example(cont.)

```
void foo_hold(struct foo *fp)                /* add a reference to the object */
{
    pthread_mutex_lock(&fp->f_lock);
    fp->f_count++;
    pthread_mutex_unlock(&fp->f_lock);
}

void foo_rele(struct foo *fp)                /* release a reference to the object */
{
    pthread_mutex_lock(&fp->f_lock);
    if (--fp->f_count == 0) {                 /* last reference */
        pthread_mutex_unlock(&fp->f_lock);
        pthread_mutex_destroy(&fp->f_lock);
        free(fp);
    } else {
        pthread_mutex_unlock(&fp->f_lock);
    }
}
```

Condition variables

Condition variables

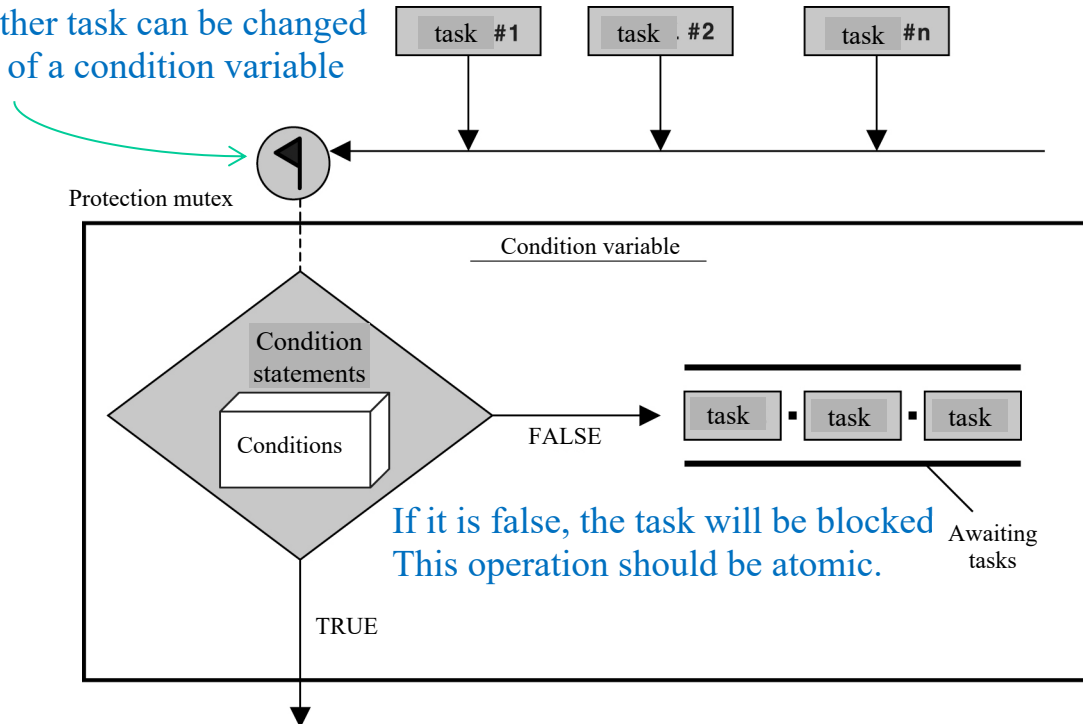
- Another synchronization mechanism
- Provides a place for threads to rendezvous.
 - When used with mutexes, it allows threads to wait for arbitrary condition to occur.
- Condition itself is protected by a mutex.
 - A thread must first lock the mutex to change the condition state.
 - Other threads will not notice the change until they acquire the mutex.

Condition variables

Condition variables

- One task can wait for other task to create a desired condition in the shared resource.

Because the status of another task can be changed while checking the status of a condition variable



Condition variables

```
#include <pthread.h>
```

```
int pthread_cond_init(pthread_cond_t *cond, pthread_condattr_t * attr);  
int pthread_cond_destroy(pthread_cond_t *cond);
```

Both return: 0 if OK, error number on failure

Pthread_cond_init

- Initializes the condition variable referenced by *cond* with attributes referenced by *attr*.
- Condition variable is represented by the `pthread_cond_t` data type.

Pthread_cond_destroy

- Destroy the given condition variable specified by *cond*.

Condition variables

```
#include <pthread.h>
```

```
int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t * mutex);  
int pthread_cond_timedwait(pthread_cond_t *cond, pthread_mutex_t *mutex,  
                           const struct timespec *timeout);
```

Both return: 0 if OK, error number on failure

Pthread_cond_wait

- Wait for a condition to be true.

Pthread_cond_timedwait

- Same with pthread_cond_wait.
- An error is returned if the specified time passes.

Condition variables

```
#include <pthread.h>
```

```
int pthread_cond_signal(pthread_cond_t *cond);
```

```
int pthread_cond_broadcast(pthread_cond_t *cond);
```

Both return: 0 if OK, error number on failure

 Notify threads that a condition has been satisfied.

 Pthread_cond_signal

- Wake up one thread waiting on a condition.

 Pthread_cond_broadcast

- Wake up all threads waiting on a condition.

Condition variables

Example

```
#include <pthread.h>

struct msg {
    struct msg *m_next;
    /* ... more stuff here ... */
};

struct msg *workq;
pthread_cond_t qready = PTHREAD_COND_INITIALIZER;
pthread_mutex_t qlock = PTHREAD_MUTEX_INITIALIZER;
```


Condition variables

Example (cont'd)

```
void process_msg(void) {
    struct msg *mp;
    for (;;) {
        pthread_mutex_lock(&qlock);
        while (workq == NULL) pthread_cond_wait(&qready, &qlock);
        mp = workq;
        workq = mp->m_next;
        pthread_mutex_unlock(&qlock);
    }
};

void enqueue_msg(struct msg *mp) {
    pthread_mutex_lock(&qlock);
    mp->m_next = workq;
    workq = mp;
    pthread_mutex_unlock(&qlock);
    pthread_cond_signal(&qready);
}
```